

Saimun Islam
CIS 4130
Saimun.Islam@baruchmail.cuny.edu

Milestone 1 Proposal:

130k Images (512x512) - Universal Image Embeddings

<https://www.kaggle.com/datasets/rhtsingh/130k-images-512x512-universal-image-embeddings>

Data Source

1. Apparel - Deep Fashion Dataset
2. Artwork - Google Scrapped
3. Cars - Stanford Cars Dataset
4. Dishes - Google Scrapped
5. Furniture - Google Scrapped
6. Illustrations - Google Scrapped
7. Landmark - Google Landmark Dataset
8. Meme - Google Scrapped
9. Packaged - Holosecta, Grozi 3.2k, Freiburg Groceries, SKU110K
10. Storefronts - Google Scrapped
11. Toys - Google Scrapped

I want to be able to predict how the future car would look like based on the dataset

I want to be able to predict how future apparel would look like based on the data set

I want to be able to predict meme and rate them based on their content from 1-10.

I want to be able to see how the landmark would look like with low and high resolution.

Milestone2 Data Acquisition:

1. To begin with I created a bucket to Amazon S3 with these code:

```
aws s3api create-bucket --bucket project-data- --region us-east-2 --create-bucket-configuration LocationConstraint=us-east-2
```

2. Install the Kaggle Command Line Interface (CLI) for Python 3.

```
$ pip3 install Kaggle
```

3. Create a new API Token in Kaggle. Go to Kaggle.com and log in. Go to the Account page.

A file kaggle.json will be downloaded.
Open up this file in Notepad or TextEdit and copy the contents to the clipboard.

4. EC2 instance:

- a. Make a directory for Kaggle: `$ mkdir .kaggle`
The leading "." is important.
- b. Create a new kaggle.json file using nano editor: `$ nano .kaggle/kaggle.json`
- c. Paste in the text from your kaggle.json file that you downloaded from Kaggle.com. It will user name:
- d. `{"username":"saimunislam1","key":"d45e468509b8f1908ead5e4b0465b4de"}`
- e. Save the kaggle.json file and exit nano.

f. Secure the file: `$ chmod 600 .kaggle/kaggle.json`
g. Try out a kaggle cli command: `$ kaggle datasets list`

Edit the `kaggle_api_extended.py` and make two changes (shown below) to accommodate writing the file to standard output.

h. Used the nano text editor to edit the `kaggle_api_extended.py` file:

`$ nano ~/.local/lib/python3.7/site-packages/kaggle/api/kaggle_api_extended.py`

i. Use the "Go To" key in Nano: `^_` to go to line 1582 in the file.

j. Change line 1582 from:

```
if not os.path.exists(outpath):  
    to
```

```
if not os.path.exists(outpath) and outpath != "-":
```

k. Change line 1594 From:

```
with open(outfile, 'wb') as out:  
    to
```

```
with open(outfile, 'wb') if outpath != "-" else  
os.fdopen(sys.stdout.fileno(), 'wb', closefd=False) as out:
```

l. Save the file and exit Nano

5. Named the data set

```
kaggle datasets download -d rhtsingh/130k-images-512x512-  
universal-image-embeddings -p - | aws s3 cp - s3://project-  
data-photos/photos.zip
```

6. S3 bucket to see if the file was downloaded

```
aws s3 ls s3://project-data-photos/
```

```
[ec2-user@ip-172-31-23-251 ~]$ aws s3 ls s3://project-data-photos/  
2022-09-30 21:16:27 13900377477 photos.zip  
[ec2-user@ip-172-31-23-251 ~]$
```

Milestone 3 Descriptive Statistics:

Step 1: I used this code to download all the data to S3

```
import zipfile
import boto3
from io import BytesIO
bucket="project-data-photos"
zipfile_to_unzip="photos.zip"
s3_client = boto3.client('s3', use_ssl=False)
s3_resource = boto3.resource('s3')
# Create a zip object that represents the zip file
zip_obj = s3_resource.Object(bucket_name=bucket,
key=zipfile_to_unzip)
buffer = BytesIO(zip_obj.get()["Body"].read())
z = zipfile.ZipFile(buffer)
# Loop through all of the files contained in the Zip archive
for filename in z.namelist():
    print('Working on ' + filename)
    # Unzip the file and write it back to S3 in the same bucket
    s3_resource.meta.client.upload_fileobj(z.open(filename
), Bucket=bucket, Key=f'{filename}')
```

step 2: I used this code to get total count of files and image

```
import boto3
import pandas as pd
s3 = boto3.resource('s3')
my_bucket = s3.Bucket('project-data-photos')
df = pd.read_csv('s3://project-data-photos/train.csv')
results = df.groupby('label').label.agg(['count'])
results.to_csv('s3://project-data-photos/train_analysis.csv')
print(results)
```

working on train.csv

```
>>>
>>> import boto3
>>> import pandas as pd
>>> s3 = boto3.resource('s3')
>>> my_bucket = s3.Bucket('project-data-photos')
>>> df = pd.read_csv('s3://project-data-photos/train.csv')
>>> results = df.groupby('label').label.agg(['count'])
>>> results.to_csv('s3://project-data-photos/train_analysis.csv')
>>> print(results)
```

	count
label	
apparel	32226
artwork	4957
cars	8144
dishes	5831
furniture	10488
illustrations	3347
landmark	33063
meme	3301
packaged	23382
storefronts	5387
toys	2402

```
>>>
```

Step3: Made a CSV file with the image details

```
from PIL import Image
import boto3
import botocore.exceptions
import csv
import pandas as pd
s3 = boto3.resource('s3')
bucket = s3.Bucket('project-data-photos')
```

```

        fields = ['Image_name', 'Width', 'Height', 'Class_name',
'Dimension']
        rows = []
        for object in bucket.objects.all():
            if 'jpg' or 'jpeg' or 'png' in object.key:
                if 'zip' in object.key:
                    continue
                print(f"Working with: {object.key}")
                name = object.key.split('/')[-1]
                slash = object.key.index('/')
                class_name = object.key[:slash]
                try:
                    bucket.download_file(object.key,
name)

                    img = Image.open(name)
                    width, height = img.size
                    dimension = f"{width}x{height}"
                    entry = [name, width, height,
class_name, dimension]

                    rows.append(entry)
                except
botocore.exceptions.DataNotFoundError as e:
                    print(e)

        df=pd.DataFrame(data = rows, columns = fields)
        df.to_csv('s3://project-data-photos/photos.csv',
index=False)

```

Step 4: I used these code to extract information about the data

```
import boto3
import pandas as pd
s3 = boto3.resource('s3')
bucket = s3.Bucket('project-data-photos')
df = pd.read_csv('s3://project-data-photos/photos.csv')
results = df.groupby('Class_name').agg({'Height':
['count', 'mean', 'min', 'max'], 'Width': ['mean', 'min', 'max']})
print(results)
```

```
>>> import pandas as pd
>>> s3 = boto3.resource('s3')
>>> bucket = s3.Bucket('project-data-photos')
>>> df = pd.read_csv('s3://project-data-photos/photos.csv')
>>> results = df.groupby('Class_name').agg({'Height': ['count', 'mean', 'min', 'max'], 'Width': ['mean', 'min', 'max']})
>>> print(results)
```

Class_name	Height				Width		
	count	mean	min	max	mean	min	max
apparel	32226	512.0	512	512	512.0	512	512
artwork	4957	512.0	512	512	512.0	512	512
cars	8144	512.0	512	512	512.0	512	512
dishes	5831	512.0	512	512	512.0	512	512
furniture	10488	512.0	512	512	512.0	512	512
illustrations	3347	512.0	512	512	512.0	512	512
landmark	33063	512.0	512	512	512.0	512	512
meme	3301	512.0	512	512	512.0	512	512
packaged	23382	512.0	512	512	512.0	512	512
storefronts	5387	512.0	512	512	512.0	512	512
toys	2402	512.0	512	512	512.0	512	512

```
>>>
```

Milestone 4 : Coding and Modeling:

```
from skimage.io import imread, imshow
import boto3
import botocore.exceptions
import csv
import numpy as np
import pandas as pd
import cv2 as cv
from PIL import ImageFile

ImageFile.LOAD_TRUNCATED_IMAGES = True
s3 = boto3.resource('s3')
bucket = s3.Bucket('project-data-photos')

fields = ['Image_name', 'Label', 'HorizontalResolution',
'VerticalResolution',\
          'LightRatio', 'RGBMean','ORB', 'FAST', 'BRIEF'] #,
'SURF', 'SIFT']
rows = []
for object in bucket.objects.all():
    if 'jpg' or 'jpeg' or 'png' in object.key:
        if 'zip' in object.key:
            continue
        print(f"Working with: {object.key}")
        name = object.key.split('/')[-1]
        slash = object.key.index('/')
        class_name = object.key[:slash]
        try:
            bucket.download_file(object.key,
name)

            # read image in RGB
            try:
                img = imread(name)
            except:
                print("Couldn't read file")
```



```

                                continue

        try:

                                width, height =
img.shape[:2]

                                except:
                                width, height = img.shape
                                pixel = np.sum(width*height)
                                rgbmean = np.average(img)
                                # Computing ORB
                                orb =
cv.ORB_create(nfeatures=50000)
                                orbkeypoints, orbdescriptor =
orb.detectAndCompute(img, None)
                                # computing FAST
                                fast =
cv.FastFeatureDetector_create()
                                fastkeypoints = fast.detect(img,
None)

                                # computing BRIEF
                                brief =
cv.xfeatures2d.BriefDescriptorExtractor_create()
                                briefkeypoints, briefdescriptor =
brief.compute(img, fastkeypoints)
                                # computing Light ratio
                                img = imread(name, as_gray=True)
                                lightratio = np.sum(img)/pixel

                                # creating entries
                                entry = [name, class_name, width,
height,lightratio, rgbmean,len(orbkeypoints),\
len(fastkeypoints),len(briefkeypoints)]
                                rows.append(entry)
                                except
botocore.exceptions.DataNotFoundError as e:
                                print(e)

                                # create .csv from dataframe
                                df=pd.DataFrame(data = rows, columns = fields)
                                df.to_csv("my_file.csv", index=False)ed

```

I use this Pyspark to the csv file and the all the images which

B2 fx apparel										
	A	B	C	D	Formula Bar	F	G	H	I	J
1	Image_name	Class	HorizontalRe	VerticalReso	LightRatio	RGBMean	ORB	FAST	BRIEF	
2	image0000.j	apparel	512	512	0.76452172	195.593187	2851	3030	2641	
3	image0001.j	apparel	512	512	0.78980799	201.652663	4920	3012	2943	
4	image0002.j	apparel	512	512	0.70704449	179.846973	1364	795	744	
5	image0003.j	apparel	512	512	0.64944974	164.162551	3277	2322	2236	
6	image0004.j	apparel	512	512	0.89960381	229.305055	2829	1342	1313	
7	image0005.j	apparel	512	512	0.75080235	191.558489	1794	936	886	
8	image0006.j	apparel	512	512	0.64923939	165.465275	2487	2403	1866	
9	image0007.j	apparel	512	512	0.63843894	165.56993	4770	5664	5448	
10	image0008.j	apparel	512	512	0.54394015	139.638008	1345	1830	1762	
11	image0009.j	apparel	512	512	0.73972266	188.511265	1785	2032	1457	
12	image0010.j	apparel	512	512	0.84622711	215.811771	1807	1217	1175	
13	image0011.j	apparel	512	512	0.55860201	143.560805	1277	1689	1471	
14	image0012.j	apparel	512	512	0.75835229	192.995941	4373	2042	1939	
15	image0013.j	apparel	512	512	0.63475857	160.667082	5627	4970	4619	
16	image0014.j	apparel	512	512	0.63711249	163.626584	1793	978	948	
17	image0015.j	apparel	512	512	0.60321394	162.216611	2193	1720	1443	
18	image0016.j	apparel	512	512	0.68855562	175.770535	1175	873	772	
19	image0017.j	apparel	512	512	0.68752571	174.641357	3145	1780	1638	
20	image0018.j	apparel	512	512	0.52574347	139.869728	3574	3384	2776	
21	image0019.j	apparel	512	512	0.91125189	232.963339	804	765	711	
22	image0020.j	apparel	512	512	0.39078091	97.4824359	4937	4180	3761	
23	image0021.j	apparel	512	512	0.67265411	170.683324	5246	6195	6105	
24	image0022.j	apparel	512	512	0.78952992	201.825232	3865	4009	3669	
25	image0023.j	apparel	512	512	0.86296381	220.343927	1979	2104	1898	
26	image0024.j	apparel	512	512	0.68147402	173.87033	602	430	388	
27	image0025.j	apparel	512	512	0.90605008	228.672808	1437	1544	1528	
28	image0026.j	apparel	512	512	0.77377634	197.696306	8451	3528	3423	
29	image0027.j	apparel	512	512	0.73041841	186.570526	3142	2904	2878	
30	image0028.j	apparel	512	512	0.59195834	151.816011	3109	3636	2975	
31	image0029.j	apparel	512	512	0.70024768	178.33307	9583	3729	3444	
32	image0030.j	apparel	512	512	0.36378861	89.4863714	4417	3924	3223	
33	image0031.j	apparel	512	512	0.78703738	200.678688	8129	3929	3765	
34	image0032.j	apparel	512	512	0.86951087	221.738747	1774	1162	1111	
35	image0033.j	apparel	512	512	0.79593489	202.752317	5360	1912	1892	
36	image0034.j	apparel	512	512	0.66729619	172.949932	9159	5373	5209	
37	image0035.j	apparel	512	512	0.69834474	178.08077	8885	4236	3858	
38	image0036.j	apparel	512	512	0.78654412	200.464031	5855	2774	2638	
39	image0037.j	apparel	512	512	0.48825736	124.978366	1997	1340	1249	

includes

Which include the file name , class, Horizontal Resolution, Vertical Resolution, RGBMean, ORB, FAST, BRIEF. After finding details about all the I wanted to predict 2 models.

The confusion matrix and Precision and Recall Curve.

To find the datasets first I created EMR 6.9.0 **Applications** Spark: Spark 3.3.0 on Hadoop 3.3.3 HDFS and Zeppelin 0.10.1. then on the EMR I downloaded this command,

```
pip3 install matplotlib
```

the code I used for modeling :

```
from pyspark.sql.functions import *
from pyspark.ml.feature import StringIndexer, OneHotEncoder,
VectorAssembler
from pyspark.ml import Pipeline
from pyspark.ml.classification import LogisticRegression,
LogisticRegressionModel
from pyspark.ml.evaluation import *
from pyspark.ml.tuning import *
from itertools import chain
from pyspark.sql.types import *
import numpy as np

spark = SparkSession.builder.getOrCreate()

photosdf = spark.read.csv("s3://project-data-
photos/my_file.csv/", header=True, inferSchema=True)

photosdf.groupby("Class").count().show()
labels={'apparel':0,
'artwork':1,'cars':2,'dishes':3,'furniture':4,
'illustrations':5, 'landmark':6, 'meme':7, 'packaged':8,
'storefronts':9, 'toys':10}
labels_expr = create_map([lit(x) for x in
chain(*labels.items())])
photosdf = photosdf.withColumn("label",
labels_expr[col("Class")])
```

```

assembler = VectorAssembler(inputCols=["HorizontalResolution",
"VerticalResolution", "LightRatio", "RGBMean", "ORB", "FAST",
"BRIEF"], outputCol="features")

photospipe = Pipeline(stages=[assembler])
transformed_photosdf =
photospipe.fit(photosdf).transform(photosdf)
x = transformed_photosdf.sort("LightRatio").sample(0.25, 1234)
x.show()
trainingData, testData =
transformed_photosdf.randomSplit([0.7,0.3])
lr = LogisticRegression(maxIter=10, regParam=0.3,
elasticNetParam=0.8)

grid = ParamGridBuilder()
grid = grid.addGrid(lr.regParam, [0.0, 0.2, 0.4, 0.6, 0.8, 1.0])
grid = grid.addGrid(lr.elasticNetParam, [0, 1])
grid = grid.build()

evaluator = MulticlassClassificationEvaluator()
cv = CrossValidator(estimator=lr, estimatorParamMaps=grid,
evaluator=evaluator)
all_models = cv.fit(trainingData)

bestModel = all_models.bestModel

test_results = bestModel.transform(testData)

test_results.select('Image_name','Class','probability',
'prediction',
'label').show(truncate=False)

print(evaluator.evaluate(test_results))

test_results.groupby('label').pivot('prediction').count().fillna
(0).show()

```

```

model_path = "s3://project-data-
photos/photos_logistic_regression"
bestModel.write().overwrite().save(model_path)

# Display Accuracy, Precision, Recall, F1 Score
def calculate_recall_precision(cm):
    tn = cm[0][1] # True Negative
    fp = cm[0][2] # False Positive
    fn = cm[1][1] # False Negative
    tp = cm[1][2] # True Positive
    precision = tp / (tp + fp)
    recall = tp / (tp + fn)
    accuracy = (tp + tn) / (tp + tn + fp + fn)
    f1_score = 2 * ((precision * recall) / (precision +
recall))
    return accuracy, precision, recall, f1_score

cm =
test_results.groupby('label').pivot('prediction').count().fillna
(0).collect()
print(calculate_recall_precision(cm))

# ROC Curve

import matplotlib.pyplot as plt

plt.figure(figsize=(5,5))

plt.plot(bestModel.summary.precisionByLabel,
bestModel.summary.recallByLabel)
plt.xlabel('precision')

plt.ylabel('Recall')

plt.title("Precision and Recall Curve")

plt.savefig("PRC1.png")

```


Milestone 5 Visualizing Results:

The confusion matrix :

label	0.0	2.0	6.0	8.0
1	435	83	708	249
6	971	368	7231	1304
3	383	11	753	593
4	2126	106	869	93
8	2125	239	996	3696
5	564	59	186	184
2	645	219	1272	257
7	432	127	198	224
0	7909	166	777	865
9	417	53	632	471
10	501	29	96	98

```
>>> # Display Accuracy, Precision, Recall, F1 Score
>>> def calculate_recall_precision(cm):
...     tn = cm[0][1] # True Negative
...     fp = cm[0][2] # False Positive
...     fn = cm[1][1] # False Negative
...     tp = cm[1][2] # True Positive
...     precision = tp / (tp + fp)
...     recall = tp / (tp + fn)
...     accuracy = (tp + tn) / (tp + tn + fp + fn)
...     f1_score = 2 * ((precision * recall) / (precision + recall))
...     return accuracy, precision, recall, f1_score
...
>>> cm = test_results.groupby('label').pivot('prediction').count().fillna(0).collect()
>>> print(calculate_recall_precision(cm))
(0.4324178782983306, 0.8159645232815964, 0.2748319641523525, 0.4111731843575419)
>>>
>>>
```

Precision and Recall Curve:



Image_name	Class	HorizontalResolution	VerticalResolution	LightRatio	RGBMean	ORB	FAST	BRIEF	label	features
image1639.jpeg	meme	512	512	0.003098877	0.794629415	210	116	116	7	[512.0,512.0,0.00...
image2425.jpeg	storefronts	512	512	0.003778379	0.9500796	642	151	151	9	[512.0,512.0,0.00...
image1605.jpeg	meme	512	512	0.022761714	5.910271962	1093	475	475	7	[512.0,512.0,0.02...
image5897.jpg	landmark	512	512	0.023430887	5.428789775	961	512	483	6	[512.0,512.0,0.02...
image3806.png	storefronts	512	512	0.030253631	9.603370667	721	311	249	9	[512.0,512.0,0.03...
image5928.png	packaged	512	512	0.03493958	8.785923004	94	54	46	8	[512.0,512.0,0.03...
image2951.png	storefronts	512	512	0.038243462	9.752082825	4640	1791	1567	9	[512.0,512.0,0.03...
image5984.png	packaged	512	512	0.039487451	8.218149821	0	1	0	8	[512.0,512.0,0.03...
image23437.jpg	landmark	512	512	0.042058982	11.00836309	1646	1216	1038	6	[512.0,512.0,0.04...
image0780.jpeg	artwork	512	512	0.045015322	12.13118617	5978	3417	3337	1	[512.0,512.0,0.04...
image0172.png	toys	512	512	0.046525385	11.9009997	2541	1309	1304	10	[512.0,512.0,0.04...
image28568.jpg	landmark	512	512	0.047250044	11.06871796	2025	1499	1423	6	[512.0,512.0,0.04...
image2136.jpeg	artwork	512	512	0.047499057	12.92299398	1333	848	841	1	[512.0,512.0,0.04...
image1586.jpeg	meme	512	512	0.05113038	12.9036115	790	394	387	7	[512.0,512.0,0.05...
image0277.png	furniture	512	512	0.053228066	13.47902552	827	320	300	4	[512.0,512.0,0.05...
image4578.png	storefronts	512	512	0.053597429	13.77123006	1761	616	586	9	[512.0,512.0,0.05...
image2274.png	furniture	512	512	0.054005601	13.51301066	223	94	94	4	[512.0,512.0,0.05...
image22881.jpg	landmark	512	512	0.056872826	13.86893972	5059	3262	3262	6	[512.0,512.0,0.05...
image5690.jpg	landmark	512	512	0.05844636	16.44370778	98	71	71	6	[512.0,512.0,0.05...
image14596.jpg	apparel	512	512	0.058646019	14.9547348	1710	712	710	0	[512.0,512.0,0.05...

only showing top 20 rows